

AGENCYOS

Master Vision & Product Specification

Document 01 of the AgencyOS Master Documentation Set

AI-FIRST AGENCY OPERATING SYSTEM

A complete operating system for running the agency lifecycle from the first WhatsApp lead through sales, onboarding, design, prototype, development, QA, payment, handover, maintenance, upsell, and repeat business — with AI agents doing the operational work and Admin controlling authority.

1. Executive Vision

AgencyOS is intended to become the agency's central operating system — not merely a CRM, project manager, invoicing tool, or chatbot.

The system coordinates the complete client journey: Facebook/other lead source → WhatsApp → AI sales → requirements → quotation → negotiation → onboarding → project → UI → prototype → development → QA → payment → handover → maintenance → upsell → repeat business.

Core vision: AI agents perform the majority of routine agency work autonomously, while human Admins retain control over authority, exceptions, commercial risk, high-impact decisions, and governance.

2. Product Definition

AgencyOS is an AI-first, workflow-driven, multi-agent agency operating system.

Workflow-driven: business events and state transitions are the foundation.

- AI-first: specialized agents perform operational work rather than merely generating suggestions.

Policy-bound autonomy: agents act independently inside approved rules.

Human-in-the-loop governance: Admin approval is required whenever an action crosses a configured authority boundary.

Auditable by design: important actions, approvals, decisions, evidence, versions, and timestamps remain traceable.

Multi-model capable: AgencyOS supports multiple AI providers and intelligent routing.

Client-transparent where appropriate: clients receive professional visibility without exposure to internal information.

3. Primary Business Objectives

- Increase lead-to-client conversion through an experienced AI Sales Agent.

Reduce manual WhatsApp communication and follow-up work.

Capture complete client requirements and convert them into structured scope.

Reduce project coordination overhead through an AI Project Manager.

Enable AI design, prototype, coding and QA agents to work continuously.

Protect the agency from payment, authorization, concurrency, and workflow errors.

Keep every important decision and action available for future audit.

Increase maintenance, upsell, retention, and repeat-project revenue.

Allow the human owner/Admin to manage exceptions instead of micromanaging routine work.

4. End-to-End Agency Journey

- Lead Arrival → WhatsApp conversation → lead record.

- AI Sales → language-aware conversation, qualification, requirement discovery and trust building.

- Quotation → AI prepares requirement-based quotation/proposal.

- Admin Approval → commercial quote/offer/exception is reviewed where required.

- Negotiation → AI handles objections and negotiates inside approved policy.

- Won → onboarding begins.

- Onboarding → client identity, email, GST preference, project data and communication context.
- Client Group → project WhatsApp group is created and connected to AgencyOS.
- Advance → invoice/payment instructions → payment proof → Admin verification.
- Project Start → required payment condition verified → project officially starts.
- UI Design → AI UI Designer creates scoped designs and themes.
- UI Approval → Admin review before client release where required; client feedback/approval captured.
- Prototype → APK/web prototype generated, tested, approved and versioned.
- Milestone → client approval → invoice → payment verification.
- Development → AI developer agents implement approved scope.
- QA → 360° testing, bug fixing, regression and production readiness.
- Final Milestone → approval → final invoice/payment verification.
- Handover → production-ready package → Admin approval → secure handover.
- Completion → client confirms completion; project history remains preserved.
- Growth → maintenance, customer success, upsell and repeat project.

5. AI Agent Operating Model

Real humans: Super Admin, Admin and Client.

Operational AI agents: Sales, Project Manager, UI Designer, Prototype, Developer, QA, Finance, Support, Customer Success, Upsell and Orchestrator/Router.

Role ≠ Actor: the architecture must support a role being performed by an AI agent, future human employee, or hybrid workflow.

- Agents may communicate and assign work to one another, with task provenance, actor, reason, inputs, outputs, approvals and audit references preserved.

6. AI Autonomy Philosophy

Autonomous does not mean uncontrolled. AgencyOS uses policy-bound autonomy.

- Agents should investigate, execute, test, retry safely, collaborate and continue until a defined completion condition is reached.
- Agents must stop at configured approval gates for commercial, financial, legal, destructive, security or production-level authority.

Every agent must have an explicit Definition of Done. Generating code or sending a message is not sufficient evidence of completion.

Failed work must trigger safe retry, alternative strategy, another specialist agent, or escalation — never silent completion.

7. Admin Authority & Approval

Admin is the authority layer, not a manual operator for every routine activity.

Admin-approved policies define what AI can do autonomously.

- AI acts automatically inside policy boundaries.
- AI requests approval when it exceeds policy or reaches a high-risk action.

Admin can approve, reject, request changes, or override where permitted.

Every approval becomes a persistent audit record linked to the exact action, artifact and version.

Typical approval-gated areas: pricing outside policy, discounts outside policy, payment-term exceptions, refunds, project cancellation, production deployment, final handover, security/permission policy changes, destructive operations and major commercial upsells.

8. WhatsApp Communication Layer

- New WhatsApp number → new lead.
- Existing completed client → new sales opportunity under the existing client.
- AI communicates in the client's language and supports text and voice-note workflows.
- Sales Agent may automatically send approved samples, portfolio, previous work, demos and follow-ups.

Quotation/offer/discount approvals can be routed to the AgencyOS internal WhatsApp group for Admin action.

A project-specific WhatsApp group is created for won clients and connected to AgencyOS.

Clients should experience professional staff communication; AI agents should not be exposed as AI agents in the client-facing identity.

WhatsApp is a communication/evidence channel, not the authoritative database. AgencyOS stores structured state, evidence, versions, approvals and audit history.

9. Sales, Pricing & Trust

The Sales Agent should operate like a highly experienced sales professional: discover the need, understand objections, build trust, explain value, and maximize conversion without inventing commitments.

Pricing is configurable. Example commercial structures such as single-vendor app/admin, website/admin, combined systems, and higher-priced multi-vendor systems are business examples, not hard-coded universal prices.

- AI may quote and negotiate automatically only inside Admin-approved pricing, discount and payment-term policies. Outside policy → Admin approval.

When a client has trust concerns, Sales Agent enters a trust-building workflow using relevant portfolio, previous work, demos, testimonials/case studies, milestone explanations and approved low-risk offers.

Approved trust options may include a smaller advance followed by UI/prototype validation, or an exceptional limited/no-advance start. AI must select from an Admin-approved playbook and must not invent commercial commitments.

10. Client Experience

The experience should feel like a professional, responsive agency team.

Language adapts to the client's language; tone remains friendly, clear and professional.

Clients receive concise explanations of what is needed, what is complete, what they must review and what happens next.

Client approvals must be explicit and easy.

Clients should not need to understand internal AI architecture.

Client-facing messages must never expose internal notes, agent conversations, credentials, private costs or internal risk data.

11. Project & Payment Philosophy

Projects are milestone-driven. A common example is 30% advance, 20% after UI/prototype, 30% after development, and 20% after testing/final delivery. Exact milestones are configurable and may be determined by AI within policy.

Core milestone rule: client milestone approval → invoice → payment → payment verification → next phase.

Where client approval is required, the client must explicitly approve the deliverable before the milestone advances.

Payment proof such as screenshot, UTR/reference or client message may create a pending verification event. Admin verification is authoritative for payment-sensitive transitions.

Finance must be transaction-safe, idempotent and auditable.

12. Design, Prototype & Development

- UI Designer Agent designs within approved scope and represents relevant scoped features and screens.

Multiple visual/theme directions may be generated.

Every client-visible design version is preserved.

Approved UI becomes the reference for prototype implementation.

Prototype work produces functional APK/web experiences that demonstrate the intended user experience.

Developer agents work from approved scope. New functionality outside scope becomes a change request.

Developer agents may code, test, debug, build and create implementation artifacts, but production-impacting actions remain policy/approval controlled.

13. Quality & Production Readiness

No agent may declare completion merely because code was generated.

- Developer Agent performs implementation and tests; QA Agent independently validates.
- Bug workflow: identify → classify severity → fix → regression → verify.

Production readiness requires objective checks such as required tests passing, P0/P1 resolution, security checks, build validation and required approvals/payment conditions.

Production deployment requires Admin approval and post-deployment smoke/live verification.

Final handover requires verified final payment, production readiness and Admin approval.

14. Multi-Model AI Strategy

AgencyOS must not be locked to one AI provider.

Supported/provider targets include Claude, OpenAI/ChatGPT models, Gemini, Grok, DeepSeek and additional providers added through a provider abstraction.

Admin Panel must contain a secure AI & Integrations page for API credentials and provider configuration.

Credentials must be encrypted/securely stored, masked in the UI, rotatable, enable/disable-able and testable.

An AI Orchestrator/Router chooses models based on configured preferences, task type, quality, speed, cost, context requirements, tool support and availability.

Fallback providers may be used when a primary provider fails.

Provider selection must remain configurable and must never expose API keys to clients or normal agent prompts.

15. Memory, Data & Audit

Every significant action should remain reconstructable months or years later.

Record actor (human/AI/system), timestamp, resource, state change, reason, policy result, approval, evidence and relevant context.

- Agent handoffs must preserve project/client context.

Projects and historical records remain preserved. Archive/deactivation may exist, but completed project history must not be casually destroyed.

Approvals must reference the exact artifact/version they approve.

16. Security & Reliability

Organization/tenant isolation must be enforced at data and authorization layers.

Critical finance operations must be transaction-safe and concurrency-safe.

Retries must be idempotent for non-repeatable operations.

Duplicate webhooks/events must not create duplicate business effects.

Errors must never silently become valid zero/empty/success states.

Secrets/API keys must never be exposed to clients.

High-risk AI permissions must be enforced by backend/application policy, not only prompts.

Critical invariants require executable tests.

17. Client Portal

Client portal should expose requirements, UI designs, prototypes, project progress, milestones, payment status, invoices, approvals, client-visible issues, deliverables and handover.

Internal AI conversations, private strategy, internal costs, credentials and security-sensitive data remain hidden.

Project progress should be derived from real tasks/milestones rather than manually fabricated.

18. Post-Completion Growth

Project completion begins customer-success operations.

- AI may identify maintenance opportunities, upsell opportunities and repeat-project opportunities.

Maintenance can renew; client payment is required before continuation unless an approved contract says otherwise.

Upsell offers use an Admin-approved catalog/policy and require approval when outside policy.

- AI may analyze objections, revisions, payment delays, support volume and conversion, but must not silently change business policy.

19. Core Product Modules

CRM/Leads; Sales AI; Quotation/Proposal; Client Onboarding; WhatsApp Operations; Project Management; Requirements/Scope; Design; Prototype; Development; QA; Finance; Client Portal; Admin Control Center; AI Orchestration; Memory/Audit; Maintenance/Upsell; Security.

20. Product Principles

- AI does the work.

Policy controls autonomy.

Approval is a real auditable state transition.

Evidence matters.

Never fake completion.

Scope is protected; new work is tracked as a change request unless clearly included.

Money is high-risk and must be safe.

Client trust is a product feature.

- AI providers are replaceable.

History is permanent.

Architecture must scale to many simultaneous leads/projects.

Future human employees and AI providers must fit the same architecture.

21. Success Criteria

From a new WhatsApp lead, AgencyOS can create the lead, communicate naturally, collect requirements, prepare a quotation, route approval, negotiate within policy, onboard the client, manage the project group, invoice and verify payments, start the project, coordinate UI/prototype/development/QA agents, obtain approvals, enforce milestone gates, prevent incomplete work from being marked complete, prepare production/handover, complete the project, preserve history, and continue into maintenance and upsell.

22. Scope of This Document

This document defines product vision, business intent, lifecycle, governance, major modules and end-state expectations.

It does not replace detailed database, API, agent, WhatsApp, finance, security, QA or deployment specifications. Those belong to the remaining AgencyOS documentation set.

If an implementation detail is not defined here, do not silently invent a business rule. Consult the relevant detailed specification or mark it Admin/Product-Decision-Required.

23. Master End-State

AgencyOS = Lead → Sales → Trust → Quote → Approval → Close → Onboard → Design → Prototype → Develop → Test → Get Paid → Handover → Complete → Maintain → Upsell → Repeat.

The final system is a highly automated agency where AI agents operate the majority of routine work continuously, specialized agents collaborate like an experienced internal team, and the human owner/Admin remains the final authority over commercial exceptions, high-risk actions, governance and strategy.

The system must be reliable enough that the owner does not repeatedly interrupt agents to ask whether work is actually complete. Completion must be proven through definitions of done, tests, reviews, evidence and appropriate gates.