

AGENCYOS

Admin Panel & Admin Control Center Specification

Document 06 of the AgencyOS Master Documentation Set

DASHBOARD • APPROVALS • MONITORING • CONFIGURATION • OVERRIDES • POLICIES •
CONTROLS

Purpose: Define the central Admin interface through which the agency owner/admin can monitor AgencyOS, approve governed actions, configure AI and business policies, intervene when necessary, manage integrations, and control the complete operation from one place.

1. Admin Control Center — Core Principle

The Admin Panel is the command center of AgencyOS. AI agents may operate autonomously, but Admin retains visibility, policy control, approval authority and emergency override capability.

- Admin should be able to manage the agency without entering the codebase.
- Routine work should remain autonomous; Admin should focus on approvals, exceptions, strategy and monitoring.
- Every important action must be auditable.
- Admin controls must be enforced by backend permissions, not merely hidden UI buttons.
- The Admin Panel must work well on desktop and mobile/responsive layouts.
- Admin should also be able to perform critical approvals through the AgencyOS internal WhatsApp group where configured.

2. Admin Dashboard

Top-level dashboard

- Total leads.
- New leads today/this week/month.
- Qualified leads.
- Active sales conversations.
- Won/lost leads.
- Active projects.
- Projects blocked.
- Projects awaiting client approval.
- Projects awaiting Admin approval.
- Pending invoices.
- Pending payment verification.
- Verified payments.
- Outstanding balances.
- Projects in UI/design.
- Projects in prototype.
- Projects in development.
- Projects in QA.
- Projects ready for handover.
- Completed projects.
- Maintenance clients.
- Upsell opportunities.
- AI agent failures.
- Provider/API failures.
- System alerts.

Dashboard behavior

- Every metric should be clickable and lead to the underlying filtered records.
- Admin can filter by date, client, project, agent, status and source.
- Dashboard should highlight exceptions rather than only showing totals.
- Critical alerts should remain visible until resolved/acknowledged according to policy.

3. Admin Navigation

- Dashboard.
- Leads & Sales.
- Clients.
- Projects.
- Approvals.
- Payments & Finance.
- Invoices.
- AI Agents.
- AI Model Routing.
- API Keys & Integrations.
- Policies & Rules.
- WhatsApp / Communication.
- Tasks & Jobs.
- QA & Releases.
- Maintenance & Support.
- Upsell / CRM.
- Analytics & Reports.
- Audit Logs.
- Security.
- System Health.
- Settings.

4. Approval Center

Approval Center is one of the most important Admin features because it prevents agents from making uncontrolled business decisions.

Approval types

- Quotation approval.
- Discount approval.
- Special offer approval.
- Out-of-policy payment plan.
- Trust-building/no-advance exception.
- Milestone completion confirmation where Admin confirmation is required.
- Payment verification.
- Refund/credit approval.
- Production deployment approval.
- Major scope-change approval.
- Client handover approval.
- Credential/security-sensitive action.
- Provider/model activation.
- AI routing policy change.
- Other configured high-risk actions.

Approval card must show

- Requesting agent.
- Client/project.
- Exact requested action.
- Reason/context.
- Amount/value where applicable.
- Relevant quotation/invoice/artifact version.
- Current policy.
- Evidence.
- Previous approval history.
- Approve / Reject / Request Changes.
- Comment/reason.
- Timestamp and Admin identity.

5. WhatsApp Admin Approval Workflow

Admin should be able to use the AgencyOS internal WhatsApp group as an operational approval channel, as previously defined.

- Agent creates structured approval request.
- Request is posted to internal Admin/AgencyOS WhatsApp group.
- Message clearly states what needs approval.
- Admin can approve/reject through the configured interaction.
- Backend validates the action and records the approval.
- Approval is linked to exact request/version.
- Agent continues only after backend confirmation.

A WhatsApp message must never be treated as an approval merely because it contains an ambiguous word such as 'okay'. The approval parser must use a controlled confirmation flow.

6. Client & Lead Management

Lead record

- Source.
- Contact information.
- Conversation history.
- Lead status.
- Sales agent assigned.
- Requirements.
- Objections/trust concerns.
- Quotation versions.
- Negotiation history.
- Follow-up schedule/status.
- Won/lost reason.

Client record

- Client profile.

- Communication preferences.
- Billing/GST details.
- All projects.
- Payments/invoices.
- Support history.
- Maintenance.
- Upsell opportunities.
- Relationship timeline.

A converted lead should become connected to the existing client record rather than producing duplicate client identities.

7. Project Control Center

Project overview

- Project name and client.
- Current lifecycle state.
- Project Manager Agent.
- Assigned agents.
- Project start date.
- Target timeline.
- Current milestone.
- Payment status.
- Scope version.
- Latest UI version.
- Latest prototype version.
- Current build.
- QA status.
- Client approval status.
- Admin approval status.
- Blockers.
- Change requests.
- Recent activity.
- WhatsApp group status.

Project actions

- Open project timeline.
- View all artifacts.
- View client feedback.
- Assign/reassign agents.
- Pause/resume project according to policy.
- Request approval.
- Add internal note.
- Create task.
- Trigger safe retry.
- Escalate to another agent.

- View payment/invoice state.
- View audit history.

8. Project Timeline & Audit View

Admin must be able to reconstruct a project's complete story from first lead contact to handover and maintenance.

- Lead creation.
- Sales conversations.
- Quotation versions.
- Admin approvals.
- Client acceptance.
- Onboarding.
- Project-group creation.
- Invoices/payment proofs.
- Payment verification.
- Project start.
- UI versions and approvals.
- Prototype versions and approvals.
- Development tasks/builds.
- QA reports.
- Final payment.
- Handover.
- Completion.
- Maintenance and upsell.

Timeline entries should show actor, timestamp, action, status and linked evidence.

9. AI Agent Control Center

Agent registry

- Agent name.
- Role.
- Status: active/paused/disabled/error.
- Current provider/model.
- Fallback provider/model.
- Current tasks.
- Queued tasks.
- Failed tasks.
- Success rate.
- Average latency.
- Usage/cost.
- Last activity.
- Permissions.

Admin actions

- Enable/disable agent.
- Pause new tasks.
- Retry failed task.
- Reassign task.
- Change model routing.
- Inspect recent executions.
- Inspect failure details.
- Update approved agent configuration.
- Force escalation to human/Admin.

10. Agent Task Monitor

- Global task queue.
- Per-agent queue.
- Running tasks.
- Waiting tasks.
- Blocked tasks.
- Failed tasks.
- Retrying tasks.
- Escalated tasks.
- Completed tasks.
- Task priority.
- Task age.
- Dependencies.
- Source/destination agent.
- Execution evidence.

Admin should be able to identify a stuck project without opening every agent's logs manually.

11. AI Model Routing Control

Provider controls

- OpenAI/ChatGPT.
- Anthropic/Claude.
- Google/Gemini.
- xAI/Grok.
- DeepSeek.
- Future providers.

Model controls

- Available/discovered models.
- Approved models.
- Blocked models.
- Experimental/canary models.
- Pinned production models.
- Preferred model per agent.

- Fallback chain.
- Capability tags.
- Cost tier.
- Quality tier.
- Latency tier.
- Health status.
- Usage.

Admin actions

- Change preferred model without code changes.
- Change fallback order.
- Disable a provider.
- Disable a model.
- Set cost limits.
- Set quality thresholds.
- Set routing policy.
- Run connection test.
- Run model test.
- View routing decisions.

12. API Keys & Integrations

All AI/API credentials should be centrally managed through an Admin-only configuration area.

- API key page in Admin Panel.
- Provider name.
- API key input.
- Masked stored key.
- Connection/test button.
- Enabled/disabled status.
- Last successful connection.
- Last error.
- Usage where available.
- Key rotation/update.
- Multiple keys per provider where supported by the architecture.
- Separate production/test credentials where appropriate.
- Never show full stored secrets after saving.
- Never send keys to client-facing agents unnecessarily.
- Encrypt secrets at rest and restrict access by role.

13. Business Policy Center

Policies should be configurable data, not scattered hard-coded conditions.

Policy categories

- Pricing.
- Discounts.

- Payment terms.
- Trust offers.
- Free vs paid change requests.
- Milestone rules.
- Approval thresholds.
- Refunds.
- Production deployment.
- Handover.
- Maintenance.
- Upsell.
- Communication.
- AI routing.
- Security.
- Agent permissions.
- Budget limits.

14. Pricing & Commercial Rules

- Configure service catalog.
- Configure product/project categories.
- Configure base pricing.
- Configure complexity tiers.
- Configure allowed price ranges.
- Configure discount limits.
- Configure special-offer rules.
- Configure payment plans.
- Configure milestone percentages.
- Configure GST/non-GST billing behavior.
- Configure account/payment-method availability.
- Configure which actions require Admin approval.

AI may recommend pricing or offers, but the Admin Policy Engine remains the authority.

15. Payment & Finance Control

Finance dashboard

- Pending invoices.
- Paid invoices.
- Overdue invoices.
- Payment proofs awaiting verification.
- Verified payments.
- Mismatched payments.
- Outstanding balances.
- Refund/credit requests.
- Milestone payment status.
- Final payment status.

Payment verification

- Admin sees expected amount, claimed amount, proof/reference and invoice.
- Approve verified payment.
- Reject/mismatch.
- Request additional evidence.
- Add internal reason/comment.
- Every verification action is audited.

Payment VERIFIED must be a backend state transition, not a free-text note.

16. Client Approval Control

- UI approval.
- Prototype approval.
- Development milestone approval.
- Final acceptance.
- Change-request approval.
- Handover confirmation.
- Approval version.
- Approval timestamp.
- Client identity.
- Approval evidence.

Admin can inspect exactly what the client approved, not merely that 'client approved' exists.

17. Scope & Change Request Control

- View latest approved scope.
- View scope history.
- Compare versions.
- Mark feature in-scope/out-of-scope.
- Approve/reject change request.
- Set free/paid classification according to policy.
- Create revised quote when required.
- Link change request to project milestone.
- Prevent developers from treating rejected/out-of-scope requests as approved work.

18. WhatsApp & Communication Control

- WhatsApp connection status.
- Connected number/account.
- Project groups.
- Internal Admin group.
- Message delivery status.
- Failed messages.
- Pending messages.
- Templates.

- Language settings.
- Automation rules.
- Approval-message formats.
- Client communication logs.
- Agent communication logs.
- Manual send where authorized.
- Emergency pause of automated messaging.

Admin should be able to stop automated messaging quickly if a provider or agent begins behaving incorrectly.

19. Automation & Workflow Control

- View active automations.
- Enable/disable workflow.
- Pause a specific project automation.
- Retry failed automation.
- View workflow dependencies.
- Set timeouts.
- Set retry limits.
- Configure escalation rules.
- Configure reminder rules.
- Configure client follow-up rules.
- Configure maintenance/upsell campaigns.

Automation changes should be versioned and audited.

20. System Health Center

Monitor

- Application health.
- Database health.
- Queue/job health.
- WhatsApp integration health.
- Email integration health.
- AI provider health.
- API error rate.
- Webhook health.
- Storage health.
- Background worker health.
- Deployment/version.
- Recent incidents.

Alerts

- Critical integration failure.
- Payment verification failure.
- AI provider outage.
- Queue backlog.

- Repeated agent failure.
- Database/storage issue.
- Security alert.
- Budget threshold exceeded.

21. AI Usage & Cost Dashboard

- Total AI usage.
- Cost by provider.
- Cost by model.
- Cost by agent.
- Cost by project.
- Cost by client.
- Cost by day/week/month.
- Input/output usage where available.
- Fallback cost.
- Highest-cost workflows.
- Budget consumption.
- Projected usage.
- Budget alerts.

This enables Admin to decide whether a premium model is being used where a cheaper approved model would be sufficient.

22. Quality & Agent Performance Dashboard

- Agent success rate.
- Task completion rate.
- Failed task rate.
- Retry rate.
- Escalation rate.
- Average completion time.
- Client-facing error rate.
- QA defect rate.
- Model/provider quality score.
- Fallback frequency.
- Human intervention frequency.
- Incomplete-deliverable detection rate.

Performance analytics should help identify agents that appear active but are not actually producing reliable work.

23. Override Center

Overrides are powerful and must be treated as exceptional actions.

Possible overrides

- Pause agent.
- Resume agent.
- Reassign task.

- Change workflow state where authorized.
- Force retry.
- Cancel task.
- Disable provider.
- Change model.
- Pause client messaging.
- Pause project automation.
- Approve an exceptional action.
- Correct an erroneous business record through controlled correction workflow.

Override safeguards

- Require confirmation for high-risk overrides.
- Require reason/comment.
- Record Admin identity.
- Record before/after state.
- Record timestamp.
- Record linked request/evidence.
- Never allow a UI override to bypass backend authorization.

24. Emergency Kill Switches

- Disable all outbound AI messaging.
- Disable one agent.
- Disable one provider.
- Disable automated payments/reminders.
- Pause project automations.
- Pause new project starts.
- Disable production deployment actions.
- Switch AgencyOS to degraded/manual mode.

Emergency controls must be reachable quickly and must not depend on the failing subsystem they are intended to stop.

25. Security & Access Control

Admin roles

- Super Admin — full control.
- Admin — operational control according to assigned permissions.
- Finance Admin — finance/payment scope.
- Operations Admin — projects/agents/workflows.
- Support Admin — client support scope.
- Read-only Auditor — view-only access.

Exact roles are configurable.

Security

- MFA where supported.
- Session/device management.

- Role-based access control.
- Least privilege.
- Sensitive action confirmation.
- Audit logging.
- Secret masking.
- IP/device controls where required.
- Security event alerts.

26. Audit Log

Anything that can materially affect money, client commitments, project state, AI behavior or security must be auditable.

- Who performed action.
- Whether human or AI.
- Agent identity.
- Provider/model where relevant.
- Action.
- Before state.
- After state.
- Timestamp.
- Reason/comment.
- Policy version.
- Approval reference.
- IP/session/device metadata where appropriate.
- Linked project/client/task.
- Evidence/artifact reference.

Audit logs should be protected from ordinary modification. Corrections should create additional audit events rather than silently rewriting history.

27. Knowledge & Configuration Center

- Service catalog.
- FAQs.
- Sales scripts/playbooks.
- Trust-building material.
- Portfolio references.
- Project templates.
- QA checklists.
- Handover checklists.
- Maintenance plans.
- Upsell catalog.
- Communication templates.
- Policy documents.
- Agent instructions/configuration.

Knowledge changes should have versions and effective dates so the Admin can understand which knowledge an agent used.

28. Agent Prompt/Configuration Management

- View agent configuration.
- Edit approved system instructions where authorized.
- Version prompts/configuration.
- Test changes in a sandbox/canary mode.
- Compare versions.
- Roll back.
- Assign configuration to agent.
- Record who changed it and why.

Prompt changes are production configuration changes and must be treated with the same discipline as software configuration.

29. Project Assignment & Workforce Control

- Assign Project Manager Agent.
- Assign UI Agent.
- Assign prototype agent.
- Assign developer agents.
- Assign QA.
- Assign support/customer-success ownership.
- Set agent priority.
- Set task concurrency where configured.
- Pause an agent for a project without disabling it globally.
- Reassign tasks when an agent/provider fails.

30. Admin Notifications

- New approval request.
- Payment verification request.
- Out-of-policy discount.
- High-value quote.
- Project blocked.
- Repeated agent failure.
- Provider outage.
- Critical QA failure.
- Security alert.
- Budget threshold.
- Final handover ready.
- Client escalation.

Notifications should support in-app alerts and the configured internal WhatsApp Admin group.

31. Reporting

- Sales conversion report.
- Lead source performance.
- Revenue report.
- Outstanding payments.
- Project delivery report.
- Milestone completion report.
- Agent productivity report.
- AI cost report.
- Provider performance report.
- QA defect report.
- Client retention report.
- Maintenance revenue.
- Upsell revenue.
- Repeat-client report.
- Exception/override report.

32. Admin Search

Admin should be able to search the entire operating system quickly.

- Search by client.
- Phone/email.
- Lead.
- Project.
- Invoice.
- Payment reference.
- Task.
- Agent.
- Approval.
- Artifact/version.
- Audit event.
- WhatsApp group.
- Search results must respect Admin's permissions.

33. Admin UX Requirements

- Clear status badges.
- Strong distinction between pending, verified, approved, rejected, blocked and completed.
- Confirmation dialogs for high-risk actions.
- Bulk operations only for safe actions.
- Filters and saved views.
- Responsive/mobile-friendly layout.
- Fast access to approvals and alerts.
- Clickable metrics.
- Readable audit timeline.

- Direct links between client, project, invoice, task and agent.
- No hidden critical status.
- No destructive action without clear warning.

34. Backend Enforcement

The Admin Panel is a UI over an authorization system; it is not the authorization system itself.

- Every Admin action must call a backend-authorized command.
- Backend checks role and permission.
- Backend checks workflow state.
- Backend checks policy.
- Backend checks required approvals.
- Backend validates data invariants.
- Backend writes audit event.
- Only then is state changed.

This prevents an agent, browser request or manipulated UI from bypassing Admin controls.

35. No Silent Overrides

- Admin override must always be visible in project history.
- AI agents must know when a human override changed their workflow state.
- A later automation must not accidentally undo an intentional Admin override.
- Overrides should carry expiry or scope where appropriate.
- Temporary emergency overrides should automatically return to normal state only when explicitly configured.

36. Acceptance Criteria

- Admin can see the complete agency dashboard.
- Admin can inspect every active project.
- Admin can approve/reject governed requests.
- Admin can verify payments.
- Admin can view exact client/project/artifact context before approval.
- Admin can manage AI providers and API keys.
- Admin can configure agent-to-model routing.
- Admin can enable/disable agents.
- Admin can pause automated messaging.
- Admin can manage pricing and business policies.
- Admin can inspect agent tasks and failures.
- Admin can override supported workflow actions.
- Every high-risk action is audited.
- Admin can access the system from mobile/responsive UI.
- Emergency kill switches are available.
- Backend authorization prevents UI bypass.
- Admin can search across clients/projects/invoices/tasks/audit records.
- System remains operational in degraded AI mode.

37. Final Admin Control Architecture

ADMIN DASHBOARD → APPROVAL CENTER → CLIENTS/LEADS → PROJECTS → AGENTS →
AI ROUTING → API KEYS → POLICIES → FINANCE → COMMUNICATION → AUTOMATION →
SYSTEM HEALTH → AUDIT → SECURITY

The Admin Panel is the control center that makes the entire AgencyOS workforce governable. AI agents can work continuously, but Admin always has the visibility and authority required to protect clients, money, projects, security and the agency.

The goal is not to make Admin micromanage AI. The goal is to make the system autonomous enough to work independently while giving Admin a clear, fast and reliable way to approve, monitor, intervene and recover.